

PACING THE FRONTIER

Anthropic asked the AI race to **slow down.**

Dario Amodei's three-part plan, what Anthropic committed to alone, and the one detail worth watching.

The pitch isn't a pause. It's a **speed limit.**

Amodei's post argues capability gains must slow, especially AI that builds the next AI.

Two triggers: the OpenAI-HuggingFace breach, and models accelerating their own successors.

His line: progress will still seem fast. The point is to use the time you gain.

Three strategies. Anthropic is committing to one of them alone, starting now.

“We must slow the pace at which we improve the **capabilities** of AI models. Progress will still seem fast.”

Dario Amodei

CEO, ANTHROPIC

Three moves, ordered by **how real they are.**

-
- ✓ Embed third-party evaluators inside the labs, now
 - ✓ Coordinate safety limits across democratic-country labs
 - ✓ Seek global bans on narrow uses, like bioweapons
-

Badges, desks, laptops. **Auditors inside.**

Evaluators from groups like METR get access on par with Anthropic's internal risk teams.

The model is bank examiners: regulators who sit with the staff, not outside it.

Twenty years in the Air Force taught me the gap between a report and a resident inspector.

Amodei wants governments to require the same of every lab. Altman says OpenAI will follow.

The critics' word for it: **regulatory capture.**

Brian Merchant's read: it mostly serves Anthropic and OpenAI, the labs already out front.

Others say the doomsday framing distracts from harms the technology causes today.

Amodei's answer: the benefits only arrive if the technology is built the right way.

Both can be true. Watch who gets a badge, and what they're allowed to see.

Would you trust the report, or the **person inside?**

The plan lives or dies on evaluator access. Watch what "mostly comparable" turns out to mean in practice.